

EXPLAINABLE AI COMPONENTS FOR NARRATIVE MAP EXTRACTION



Authors

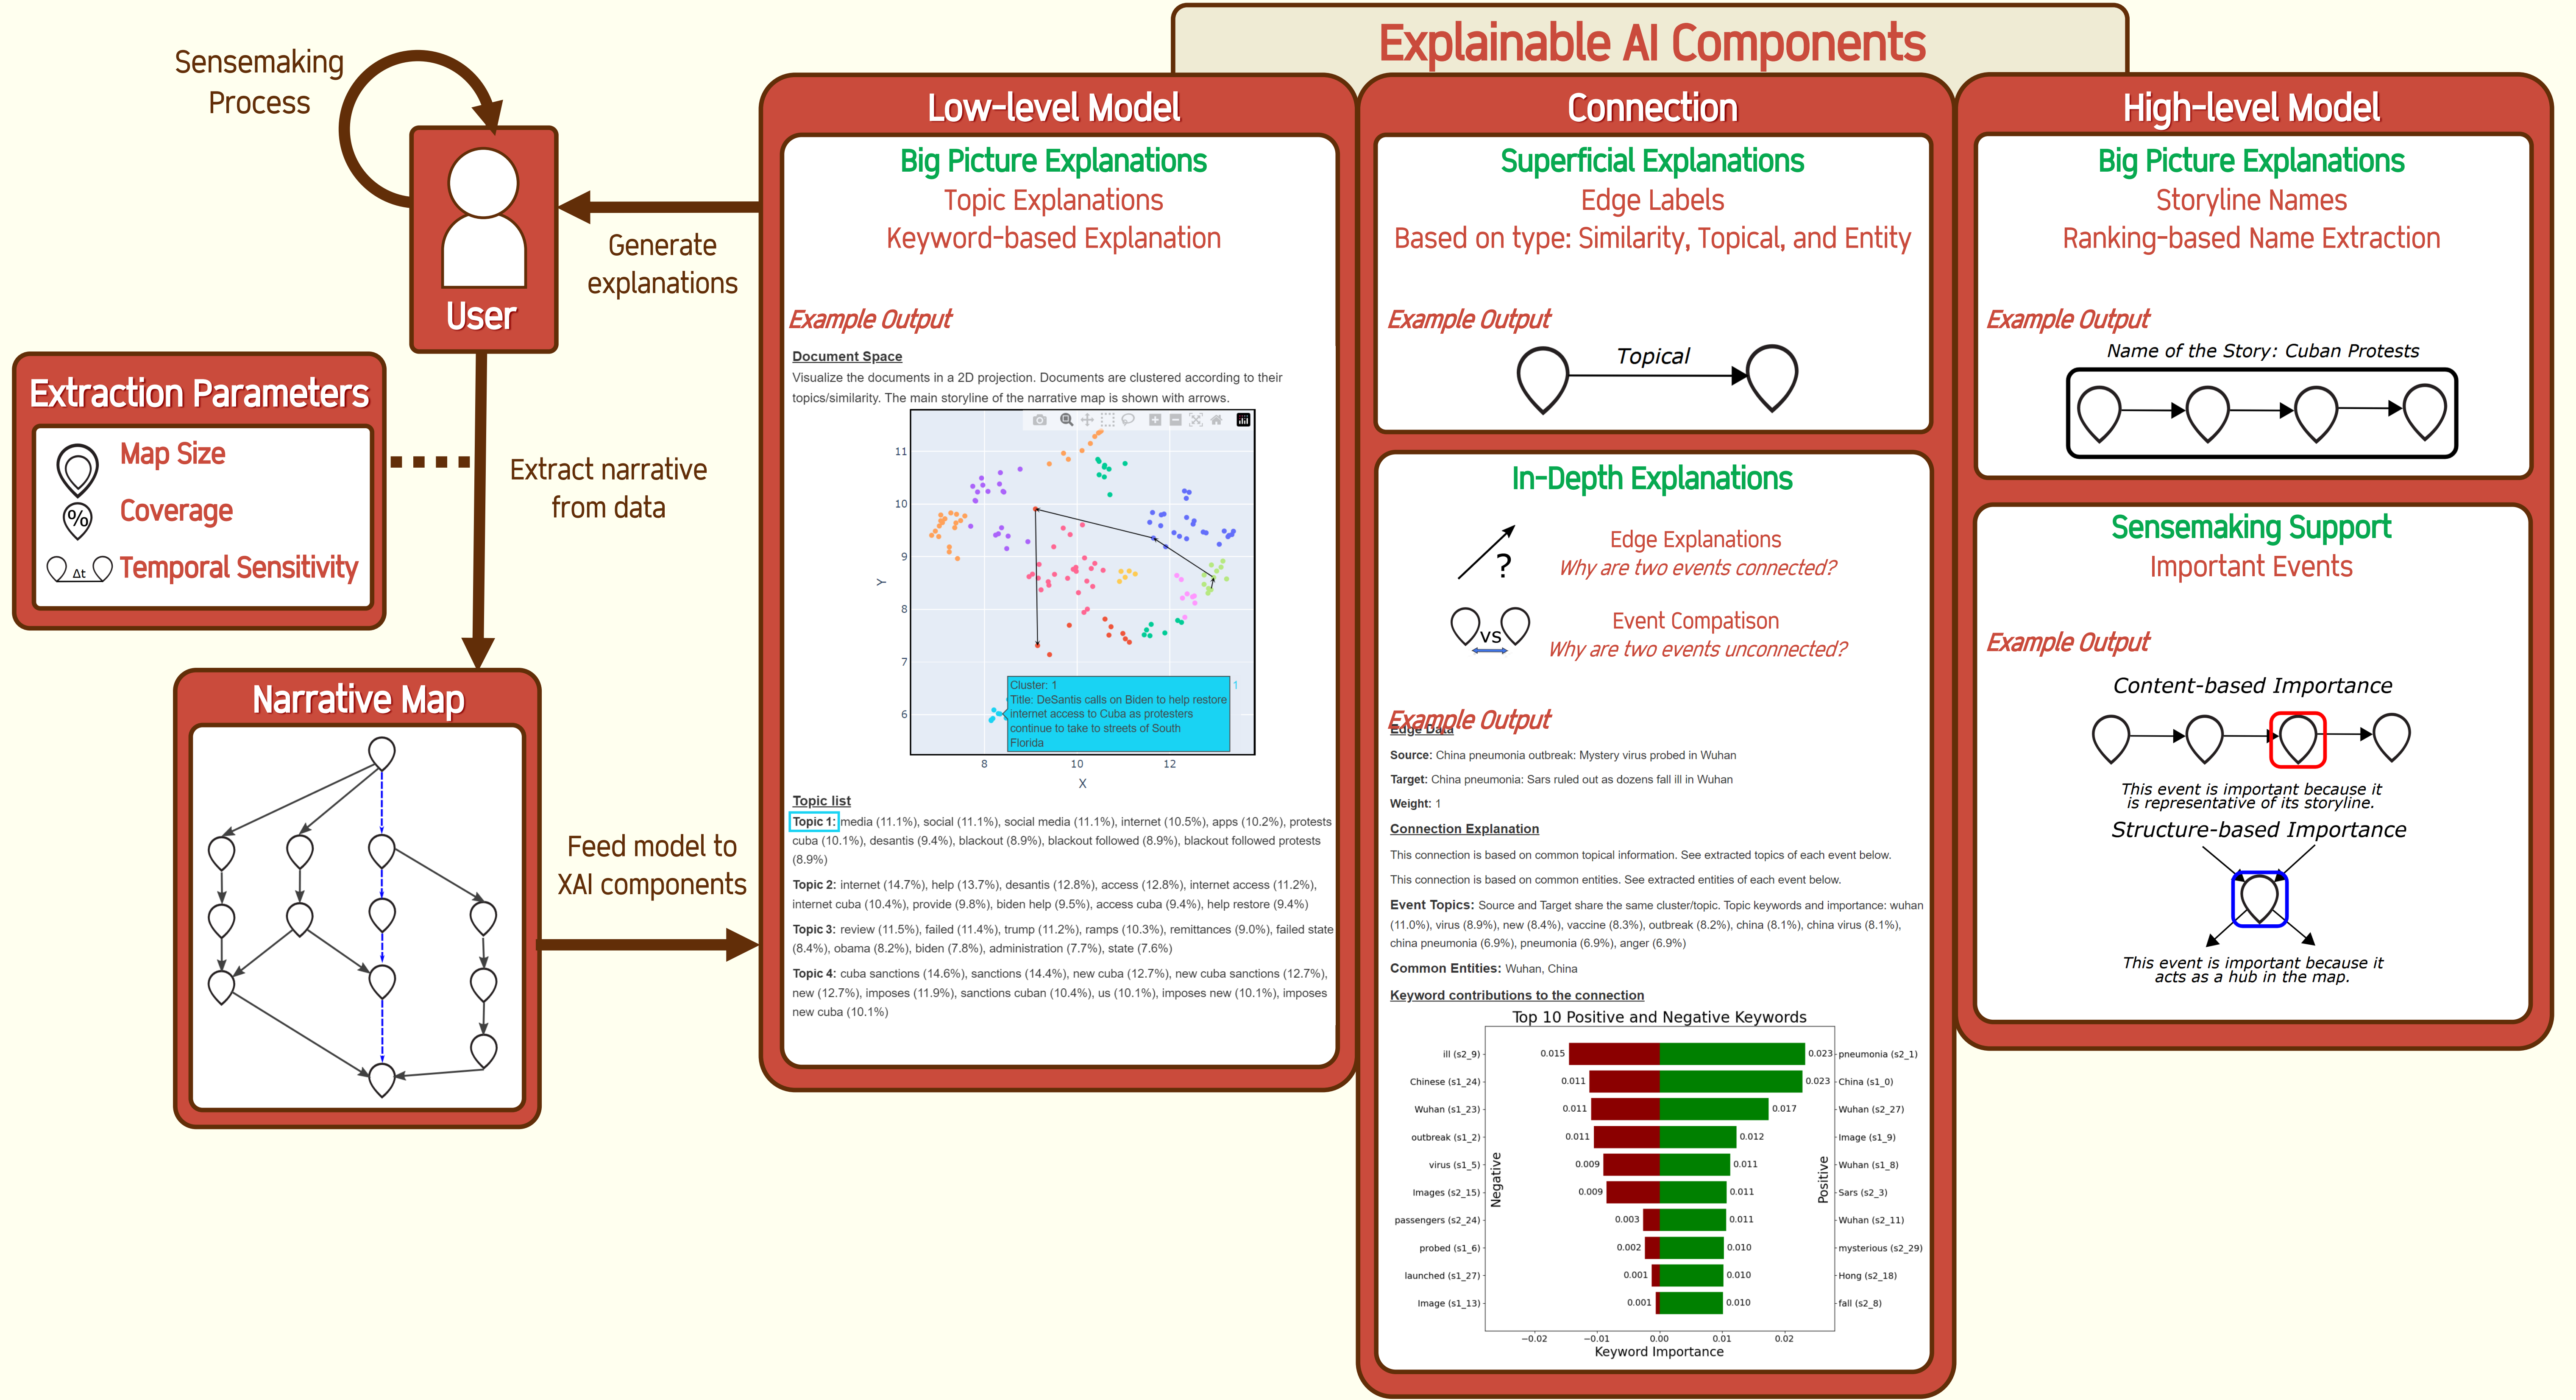
Brian Keith, Fausto German, Eric Krokos, Sarah Joseph, Chris North
Universidad Católica del Norte, Chile | Virginia Tech, USA | U.S. Government

Research Problem

- Narrative extraction systems are becoming increasingly complex.
- Users need to understand why and how AI systems make decisions.
- Trust – in the broad sense of reliability, predictability, and efficiency – requires explanations at multiple levels of abstraction.

Motivation

- Complex AI-based narrative extraction systems appear as "black boxes" to users.
- Analysts need to understand and trust why specific narrative connections were made.
- Effective explanations must connect text details with the overall narrative structure.



Overview of XAI System for Narrative Maps

Low-level Space Explanations

- Topical clusters using HDBSCAN clustering.
- Keyword-based explanations with modified TF-IDF.
- Interactive 2D visualization with UMAP projection.

Connection Explanations

- Classification of connections: similarity-based, entity-based, topical.
- SHAP value analysis to identify influential terms.
- Event comparison to explain both connected and unconnected events.

High-level Structure Explanations

- Automated storyline naming through noun phrase extraction.
- Important event detection using:
 - Content importance (similarity to storyline centroid).
 - Structural importance (degree centrality in narrative graph).

Evaluation Results

- User study with 10 participants analyzing 2021 Cuban protests narrative.
- Overall explanations increased user trust ($M = 4.5/5$).
- Important event detection was considered highly relevant ($M = 4.3/5$).
- Storyline naming showed mixed results (correctness: $M = 3.0/5$).

Key Contributions

- Multi-level explanation approach bridges document space and narrative structure.
- Connection explanations and important event detection build user confidence.
- System helps users develop appropriate trust in narrative extraction.
- Available on GitHub:** <https://github.com/briankeithn/narrative-maps>.



Future Directions

- More sophisticated temporal and causal explanation strategies.
- Adaptive explanations based on user expertise.
- Improving scalability for larger narrative collections.
- Better alignment between the model and explanations.

Responses (5-point Likert Scale)

1 ("Strongly Disagree") to 5 ("Strongly Agree")

(a) Overall Explanations							(b) Important Events						
	1	2	3	4	5	Mean SD		1	2	3	4	5	Mean SD
Trust	0	0	1	3	6	4.50 0.71	Relevance	0	0	1	5	4	4.30 0.67
Usefulness	0	0	1	5	4	4.30 0.67	Usefulness	0	0	2	4	4	4.20 0.79
(c) Storyline Names							(d) Connections						
	1	2	3	4	5	Mean SD		1	2	3	4	5	Mean SD
Correctness	1	4	1	2	2	3.00 1.41	Label Correctness	0	0	2	6	2	4.00 0.67
Relevance	1	0	4	1	4	3.70 1.34	Label Usefulness	0	1	3	5	1	3.60 0.84
Usefulness	0	1	2	5	2	3.80 0.92	Explanation Usefulness	0	0	1	7	2	4.10 0.57
							Comparison Usefulness	0	0	4	4	2	3.80 0.79